

Index

1.	Question 1.....	1
2.	Question 2.....	1
3.	Question 3.....	3
4.	Question 4.....	3
	Appendices	5
	APPENDIX A: PYTHON CODE FOR QUESTION 3.....	5
	APPENDIX B: PLOT FOR QUESTION 3.....	12
	APPENDIX C: DATA DISTRIBUTION	12

1. Question 1

In a Support Vector Machine (SVM), we aim to separate two classes using a decision boundary (hyperplane). This boundary is defined as: $w^T x_i + b = 0$, where w represents the weight vector, b is the bias that adjusts the position of the hyperplane.

Next, define two supporting hyperplane ensuring that data points fall on the correct side of the margin.

- For positive class ($y_i = +1$): $w^T x_i + b \geq +1$
- For negative class ($y_i = -1$): $w^T x_i + b \leq -1$

To simplify and unify these conditions, we multiply y_i which is $+1$ for positive class and -1 for negative class, resulting in: $y_i(w^T x_i + b) \geq 1$. This constraint guarantees that all points are at least 1 unit away from the decision boundary.

- For positive points ($y_i = +1$), it remains $w^T x_i + b \geq 1$
- For negative points ($y_i = -1$), it transforms into $w^T x_i + b \leq -1$

Finally, to maximise the margin, the margin width (distance between the two hyperplanes) is given by $\frac{2}{\|w\|}$. Since the larger margin leads to better generalization, we minimize $\|w\|$ and make it mathematically easier, thus, the final SVM optimization problem becomes: $\min_{w,b} (1/2) \|w\|^2$ subject to: $y_i(w^T x_i + b) \geq 1$, for all i

2. Question 2

2.1. Choose initial feasible solution

Given that initial feasible solution $z_0 = (w_0, b_0)$ to satisfy the constraints:

$$y_i(w_0^T x_i + b_0) \geq 1$$

$$-M \leq w_j \leq M, \quad j = 1, \dots, d$$

$$-M \leq b \leq M$$

If the dataset is linearly divisible, initialise the $w_0 = 0, b_0 = 0$, and then adjust it to satisfy the constraints.

2.2. Calculate the gradient

- Objective function: $f(z) = \frac{1}{2} \|w\|^2 = \frac{1}{2} \sum_{j=1}^d w_j^2$
- Find the gradient for w : $\nabla f(w, b) = (w, 0)$

2.3. Calculate the feasible direction $d_k = v_k - z_k$

v_k is obtained by solving the optimisation problem: $v_k = \arg \min_{v \in Z} \nabla f(z_k)^T v$, and is chosen as the feasible point in the direction that decreases the objective function the most.

Plus, objective function in SVM is: $f(w) = \frac{1}{2} \|w\|^2$, its gradient is: $\nabla f(w) = w$

Thus, we need to find the closest feasible point for minimising the gradient. The projection operator $\text{Proj}_Z w_k$ finds the closest feasible point to w_k in the feasible set $v_k = -\text{Proj}_Z(w_k)$. Projection naturally results from gradient minimisation, which results in: $d_k = -\text{Proj}_Z(w_k) - w_k$

2.4. Calculate the step τ_k

τ_k needs to be solved by the following optimisation problem: $\tau_k = \arg \min_{\tau \in [0,1]} f(zk + \tau dk)$

Then: $\tau k = \arg \min_{\tau \in [0,1]} \frac{1}{2} \|w_k + \tau d_k\|^2$

Calculate the derivative of τ_k to find the optimal solution: $\frac{1}{2} \cdot 2(w_k + \tau d_k)^T d_k = (w_k + \tau d_k)^T d_k$

Let the derivative be equal to zero: $(w_k + \tau d_k)^T d_k = 0$

Result: $\tau_k = -\frac{w_k^T d_k}{|d_k|^2}$ Ranging from $[0,1]$. Since τ is constrained to $[0,1]$, project τ^* into this range:

$$\tau^* = \max \left(0, \min \left(1, -\frac{w_k^T d_k}{|d_k|^2} \right) \right)$$

- If $\tau^* < 0$: This would imply moving in the opposite direction, which is not allowed in this optimisation context. So, we set the step size to 0 (no movement).
- If $\tau^* > 1$: This would imply taking a step larger than the optimal direction suggested by the gradient and feasible direction, which might cause us to exceed the boundaries. So, we set the step size to 1 (maximum allowed movement).
- If $0 \leq \tau^* \leq 1$: If the computed τ^* already lies in the feasible range, then we can directly use it without modification.

2.5. Updating solutions $z_{(k+1)} = z_{(k)} + \tau^* d_{(k)}$

2.6. Checking for convergence

If $|z_{k+1} - z_k| \leq \epsilon$ then stop, otherwise return to Step 2.2 and continue iteration.

3. Question 3

The algorithm halts when the change in the solution z_k , (including weight w and bias b) between iterations falls below a predefined tolerance ϵ , calculated with the Euclidean norm $\|z_k - z_{k-1}\|$. For classifier, the model perfectly distinguishes the Iris-setosa class (label = +1) and other class (label = -1), achieving 100% accuracy with the classifying weight of sepal length, sepal width, petal length, petal width shows -0.04604523, 0.5216668, -1.0032087, and -0.46445877, respectively. Plus, the bias is 1.45118752 for adjusting the position of the decision boundary. Additionally, tolerance parameter ϵ is set as 1e-5, a smaller ϵ means the algorithm will continue iterating until the change in z_k is very small, resulting in a more accurate solution. Finally, M is set as 10, this means that each component of w and b is constrained to lie within the interval $[-M, M]$. Through cross-validation, the small M value 10 restricts the solution space, gives the best performance on a validation set is selected.

4. Question 4

Knowing that Setosa and Versicolor/Virginica are linearly distinguishable, and Versicolor and Virginica are non-linearly distinguishable, therefore, the data's noise and class overlap necessitate a soft-margin SVM with a regularisation parameter C . To balance misclassification and generalisation, optimising C through cross-validation is crucial, shifting from a rigid to a more adaptable hyperplane.

4.1. Model Interpretation and Variable Assumption Method

4.1.1. Additional Decision Variables

We introduce slack variables u_i for each training sample. These variables allow the SVM to handle cases where the data is not perfectly linearly separable, where u_i measures the degree to which the training sample violates the margin constraints.

4.1.2. New Objective Function $\min_{w,b,u} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^N u_i$

$\frac{1}{2} \|w\|^2$ represents the original term that maximises the margin by minimising the norm of w .

$\sum_{i=1}^N u_i$ represents a penalty term for violating the margin constraint, where C is a hyperparameter controlling the trade-off between maximising the margin and minimising classification errors.

4.1.3. Modified Constraints

Margin Constraint: $(w^T x_i + b) \geq 1 - u_i, \forall i = 1, 2, \dots, N$. This allows for a margin violation when $u_i > 0$, enabling SVM to handle non-linearly separable data.

Non-negativity Constraint: $u_i \geq 0, \forall i = 1, 2, \dots, N$. This ensures that the slack variables are non-negative.

4.1.4. Set the Value for Parameter C

A larger C means the SVM will try to avoid misclassification more strictly, resulting in a smaller margin (risk of overfitting). A smaller C allows for a larger margin, tolerating more classification errors (better generalization).

4.2. Algorithm to Classify a New Iris Record

Apply Hard-Margin SVM to distinguish Setosa first, then perform Soft-Margin SVM to distinguish Setosa.

4.2.1. Train a Hard-Margin SVM to Separate Setosa (0) and Non-Setosa (1)

Train a Hard-Margin SVM: $\min_{w,b} \frac{1}{2} \|w\|^2$ subject to: $y_i(w^T x_i + b) \geq 1, \forall i = 1, 2, \dots, N$

If the new sample is classified as Setosa, the classification is complete. Otherwise, proceed to Step 4.2.2.

4.2.2. Train a Soft-Margin SVM to Classify Versicolor (1) vs. Virginica (2)

Train a Soft-Margin SVM: $\min_{w,b} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^N u_i$ subject to: $y_i(w^T x_i + b) \geq 1 - u_i, u_i \geq 0, \forall i = 1, 2, \dots, N$

If the sample is classified as Versicolor, assign it to that class. Otherwise, classify it as Virginica.

Appendices

APPENDIX A: PYTHON CODE FOR QUESTION 3

```
from pyomo.environ import *
from pyomo.opt import SolverStatus, TerminationCondition
from sklearn.linear_model import LinearRegression
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

# Load the data
data = pd.read_csv("Data.csv", header=None)
X = data.iloc[:, :-1].values # Features
y = data.iloc[:, -1].values # Labels

# Convert labels to binary: Iris-setosa = 1, others = -1
y = np.where(y == 'Iris-setosa', 1, -1)

# Parameters
d = X.shape[1] # Number of features
M = 10 # Bound for w and b
epsilon = 1e-5 # Convergence threshold
max_iter = 1000 # Maximum number of iterations
```

```

# Initialize w and b randomly

weights = np.random.uniform(-M, M, d)

bias = np.random.uniform(-M, M)

# Objective function: 1/2 * ||w||^2

def objective(z):

    w=z[:-1]

    return 0.5 * np.dot(w, w)

# Gradient of the objective function

def gradient(z):

    w=z[:-1]

    return np.append(w, 0) # Gradient w.r.t. w is w, w.r.t. b is 0

# Solve linear programming problem to find v_k

def find_vk(z):

    model = ConcreteModel()

    # Decision variables: v = [w_v; b_v]

    model.weights_v = Var(range(d), bounds=(-M, M))

    model.bias_v = Var(bounds=(-M, M))

    # Objective: minimize gradient(z)^T v

    def objective_function(m):

```

```

        return sum(gradient(z)[j] * m.weights_v[j] for j in range(d))
+ gradient(z)[-1] * m.bias_v

model.objective = Objective(rule=objective_function,
sense=minimize)

# Constraints: y_i (w_v^T x_i + b_v) >= 1 for all i

def constraint_function(m, i):

    return y[i] * (sum(m.weights_v[j] * X[i, j] for j in range(d))
+ m.bias_v) >= 1

model.constraints = Constraint(range(len(y)),
rule=constraint_function)

# Solve the problem

solver = SolverFactory('glpk')

results = solver.solve(model)

# Extract the solution

weights_v = np.array([model.weights_v[j]() for j in range(d)])
bias_v = model.bias_v()

return np.append(weights_v, bias_v)

# compute_tau_k using IPOPT

def compute_tau(z, d):

```

```

model = ConcreteModel()

# Decision variable: tau_k
model.tau = Var(bounds=(0, 1))

# Objective: minimize the objective function at z + tau * d
def objective_function(m):
    return objective(z + m.tau * d)

model.objective = Objective(rule=objective_function,
sense=minimize)

# Solve the problem
solver = SolverFactory('ipopt')    # Use IPOPT solver
results = solver.solve(model)

# Extract the solution
return model.tau()

# Optimization loop
z = np.append(weights, bias)

z_history = pd.DataFrame(columns=[f'weight_{j+1}' for j in range(d)] +
['bias'])

z_history.loc[0] = z    # Add initial solution to the DataFrame

```

```

for iteration in range(max_iter):

    # Find v_k
    v_k = find_vk(z)

    # Find direction d_k = v_k - z
    d_k = v_k - z

    # Compute tau_k
    tau_k = compute_tau(z, d_k)

    # Update z
    z_new = z + tau_k * d_k
    z_history.loc[iteration + 1] = z_new

    # Check for convergence
    if np.linalg.norm(z_new - z) < epsilon:
        break

    z = z_new

# Extract optimal weights and bias
optimal_weights = z[:-1]
optimal_bias = z[-1]

```

```

# Predict function

def predict(X):
    return np.sign(np.dot(X, optimal_weights) + optimal_bias)

# Evaluate accuracy

y_pred = predict(X)

accuracy = np.mean(y_pred == y)

print(f"Final Accuracy: {accuracy:.4f}")

print(f"Final SVM Weights: {z}")

print(f"Final Bias: {optimal_bias}")

df_columns = ['sepal_length', 'sepal_width', 'petal_length',
'petal_width', 'class']

feature_importance = np.abs(optimal_weights)

top_2_indices = np.argsort(feature_importance)[-2:] # Get two largest
weights

print(f"Selected Features: {df_columns[top_2_indices[0]]} ,
{df_columns[top_2_indices[1]]}")

def plot_decision_boundary(X, y, w, b, feature_indices):
    plt.figure(figsize=(8, 6))

    # Use only the two selected features

```

```

X_vis = X[:, feature_indices]

w_vis = w[feature_indices] # Use the weights corresponding to
selected features

for label, color in zip([1, -1], ['#3d5a80', '#ee6c4d']):

    subset = X_vis[y == label]

    plt.scatter(subset[:, 0], subset[:, 1], label="Iris-setosa" if
label == 1 else "Others",
                color=color, alpha=0.7, edgecolors='k')

# Compute the decision boundary

x_min, x_max = X_vis[:, 0].min() - 1, X_vis[:, 0].max() + 1
x_vals = np.linspace(x_min, x_max, 100)
y_vals = -(w_vis[0] * x_vals + b) / w_vis[1]

plt.plot(x_vals, y_vals, 'k-', linewidth=2, label='Decision
Boundary')

plt.xlabel(f'Feature {df_columns[feature_indices[0]]}')
plt.ylabel(f'Feature {df_columns[feature_indices[1]]}')
plt.title("Optimized SVM Decision Boundary")

plt.legend()

plt.show()

# Select the best two features and plot

plot_decision_boundary(X, y, optimal_weights, optimal_bias,
top_2_indices)

```

APPENDIX B: PLOT FOR QUESTION 3

Bound ($M=10$): Prevents weights and bias from diverging while allowing sufficient flexibility.

Tolerance $\epsilon=1e-5$: Ensures convergence without unnecessary computations.

Max Iterations max_iter=1000: Provides enough updates for stability without excessive computation.

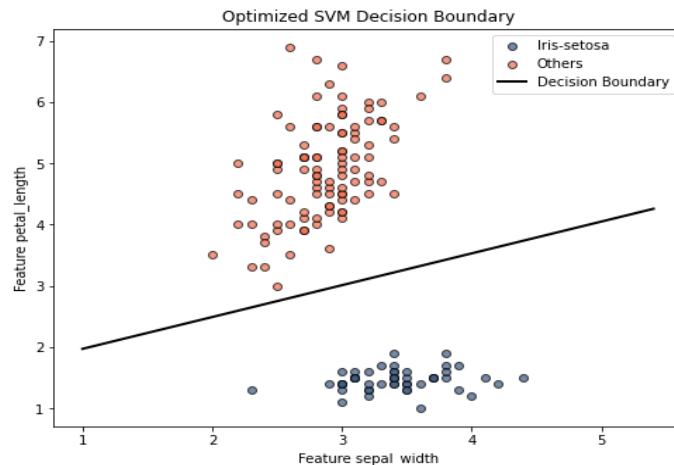


Figure 1: Sepal Width vs Petal Length Features SVM.

APPENDIX C: DATA DISTRIBUTION

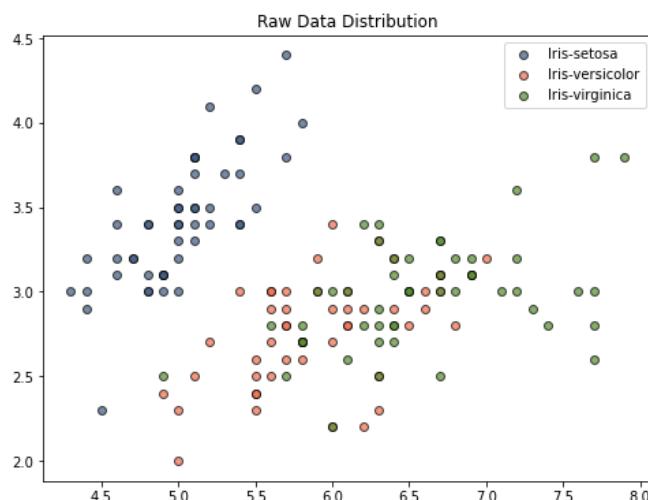


Figure 2: Data Distribution